

# Practical Work 3: Clustering, Unsupervised Analysis, and MLOps

Louis Fippo Fitime, Claude Tinku, Kerolle Sonfack  
Computer Engineering Department, ENSPY

October 11, 2025

## TP Objective

- Implement and compare the three major families of Clustering algorithms: **K-Means**, **GMM/EM**, and **DBSCAN**.
- Master unsupervised evaluation metrics (**Inertia**, **Silhouette Coefficient**).
- Study the role of **Clustering** in **Customer Knowledge (Sales)** and the **Gaming Industry**.
- Consolidate the **MLOps** phase by tracking clustering configurations (number of clusters, hyperparameters) via **MLflow**.

---

## 1 Phase I: Representative-Based Clustering (K-Means)

### 1.1 Use Case 1: Customer Knowledge and Segmentation (Sales)

We will segment customers based on their annual spending and spending score. The objective is to find homogeneous groups for targeted marketing campaigns.

- **Task 1.1: Data Preparation:** Load the customer dataset (simulation), reduce the dimension to two variables for visualization (*AnnualIncome* and *SpendingScore*)

```
Code Python 1.1: Data Preparation (Simulation)
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

Simulated data for customer segmentation (2D for visualization)

```

columns: CustomerID, Genre, Age, Annual_Income (k$), Spending_Score (1-100)
X_sim = np.array([
[15, 39], [17, 35], [18, 92], [20, 95], [23, 75], [78, 7], [75, 12],
[77, 20], [79, 3], [40, 40], [42, 52], [45, 35], [55, 60], [58, 58]
])

Data Standardization
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_sim)

print("Standardized data (excerpt):")
print(X_scaled[:5])

```

## 1.2 K-Means Optimization and Evaluation

- **Task 1.2: Choosing K (Elbow Method):** Apply the K-Means algorithm for K ranging from 1 to 10. Calculate and visualize the inertia (sum of intra-cluster squares) as a function of K. Deduce the optimal number of clusters.
- **Task 1.3: Silhouette Evaluation:** Calculate the **Silhouette Coefficient** for the chosen optimal K. Briefly recall the meaning of this metric.

```

Code Python 1.2 & 1.3: K Optimization
from sklearn.metrics import silhouette_score

inertia = []
K_range = range(1, 11)

for k in K_range:

TO BE COMPLETED: Train KMeans and store the inertia in the 'inertia' list
Elbow Method Visualization
plt.figure(figsize=(8, 5))
plt.plot(K_range, inertia, 'bo-')
plt.title("Elbow Method for Inertia")
plt.xlabel("Number of Clusters (K)")
plt.ylabel("Inertia")
plt.show()

TO BE COMPLETED: Choose K_optimal and calculate the silhouette score
K_optimal = ...
kmeans_final = KMeans(n_clusters=K_optimal, random_state=42)
labels = kmeans_final.fit_predict(X_scaled)
silhouette_avg = silhouette_score(X_scaled, labels)

print(f"\nSilhouette Coefficient for K={K_optimal}: {silhouette_avg:.4f}")

```

---

## 2 Phase II: Advanced Clustering and Critical Comparison

### 2.1 Use Case 2: User Behavior Segmentation (Gaming)

Real video game user data is often noisy and non-globular in shape. We will use synthetic data to compare models.

- **Task 2.1: Probabilistic Clustering (GMM/EM):** Explain the role of the **Expectation-Maximization (EM) Algorithm** in fitting a **Mixture of Gaussian Models (GMM)**. Train a GMM and visually compare the results with K-Means on non-globular data (e.g., two moons).

Code Python 2.1: GMM and Comparison

```
from sklearn.mixture import GaussianMixture
from sklearn.datasets import make_moons # Non-globular data
```

Generating 'moons' data

```
X_moons, y_moons = make_moons(n_samples=200, noise=0.05, random_state=42)
X_moons_scaled = StandardScaler().fit_transform(X_moons)
```

K-Means on Moons

```
kmeans_moons = KMeans(n_clusters=2, random_state=42)
labels_kmeans = kmeans_moons.fit_predict(X_moons_scaled)
```

GMM on Moons

```
gmm = GaussianMixture(n_components=2, random_state=42)
labels_gmm = gmm.fit_predict(X_moons_scaled)
```

TO BE COMPLETED: Visualize the results of the two algorithms

Use plt.scatter(X\_moons\_scaled[:, 0], X\_moons\_scaled[:, 1], c=labels\_kmeans/gmm, cmap=

### 2.2 Density-Based Clustering (DBSCAN)

- **Task 2.2: DBSCAN:** Explain the functioning of the **DBSCAN** algorithm and the interpretation of its hyperparameters (`, minPts`). Apply DBSCAN to the `X_moons_scaled` data and comment on its results.

Code Python 2.2: DBSCAN

```
from sklearn.cluster import DBSCAN
import numpy as np
```

TO BE COMPLETED: Choose epsilon and min\_samples hyperparameters

```
dbscan = DBSCAN(eps=..., min_samples=...)
labels_dbscan = dbscan.fit_predict(X_moons_scaled)
```

Display the number of clusters found (excluding noise points, label -1)

```
num_clusters = len(set(labels_dbscan)) - (1 if -1 in labels_dbscan else 0)
```

```
print(f"\nDBSCAN: Number of clusters found: {num_clusters}")
```

```
print(f"Number of noise points (label -1): {np.sum(labels_dbscan == -1)}")
```

## 2.3 Neighbor and Similarity Concepts

- **Task 2.3: k-NN and Metrics:** Describe the difference between the objectives of **Supervised Learning** (k-NN in classification/regression) and **Unsupervised Learning** (Clustering). Cite at least three relevant similarity metrics for textual or high-dimensional data, beyond Euclidean distance (e.g., Cosine, Jaccard).
  - **Task 2.4 (Bonus): Search Optimization:** Explain how **KD-Trees** and **Locality-Sensitive Hashing (LSH)** structures are used to reduce the time complexity of nearest neighbor search in large databases (e.g., a document-based recommendation system).
- 

## 3 Phase III: MLOps and Clustering Model Management

As seen in the previous TP, traceability is essential. Choosing the optimal number of clusters ( $K$ ) is a crucial hyperparameter.

### 3.1 Segmentation Model Tracking with MLflow

- **Task 3.1: Experiment Structuring:** Using the **MLflow** approach, structure the code to log the results of the K-Means experiment (Phase I). The objective is to later compare the impact of the choice of  $K$  (hyperparameter) on Inertia and the Silhouette Score (metrics).

```
Code Python 3.1: Clustering Configuration Tracking
import mlflow
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, calinski_harabasz_score
import joblib

TO BE COMPLETED: Ensure MLflow is configured (mlflow.set_experiment)
def track_clustering_run(X_data, k, run_name):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans.fit_predict(X_data)

    Calculate metrics
    score_silhouette = silhouette_score(X_data, labels)
    inertia_value = kmeans.inertia_

    with mlflow.start_run(run_name=run_name):
        # 1. Logging Hyperparameters
        mlflow.log_param("K", k)
        mlflow.log_param("Algorithm", "KMeans")

        # 2. Logging Metrics
        mlflow.log_metric("Inertia", inertia_value)
        # TO BE COMPLETED: Log the Silhouette Coefficient
```

```
# 3. Logging Artifact (Centroids)
# joblib.dump(kmeans, "kmeans_model.pkl")
# mlflow.log_artifact("kmeans_model.pkl")
```

Example tracing for K=3 and K=4

```
Assuming X_scaled is available from Phase I
track_clustering_run(X_scaled, 3, "KMeans_K3_Initial")
track_clustering_run(X_scaled, 4, "KMeans_K4_Alternative")
print("\nMLflow tracking of K-Means experiments finished.")
```

## 3.2 Segmentation Model Deployment

- **Task 3.2: Packaging (Recap):** Once the optimal K is chosen, explain how the selected clustering model (K-Means) and the associated `StandardScaler` would be **packaged** (e.g., via `joblib`) for integration into a Dockerized web API (recap from previous TP). What is the role of this package in production for the marketing department?
  - **Task 3.3: Post-Production (Monitoring):** If this customer segmentation model is deployed, cite and briefly describe two specific indicators of **Data Drift** or **Model Decay** that you should monitor (e.g., centroid drift, cluster size distribution).
- 

## 4 Advanced Theoretical Concepts (EM, LDA)

- **Task 4.1: Initialization Comparison:** Discuss the problems of the non-convex optimization objective in Clustering (especially K-Means and GMM). Cite and compare two initialization techniques (e.g., K-Means++ vs. Random Initialization) to better converge toward a global optimum.
- **Task 4.2: Mixed Membership Modeling (LDA):** Describe the concept of **Latent Dirichlet Allocation (LDA)** for Topic Modeling in documents. How does LDA differ from K-Means clustering applied directly to word vectors (TF-IDF), particularly through the idea of *mixed membership*?